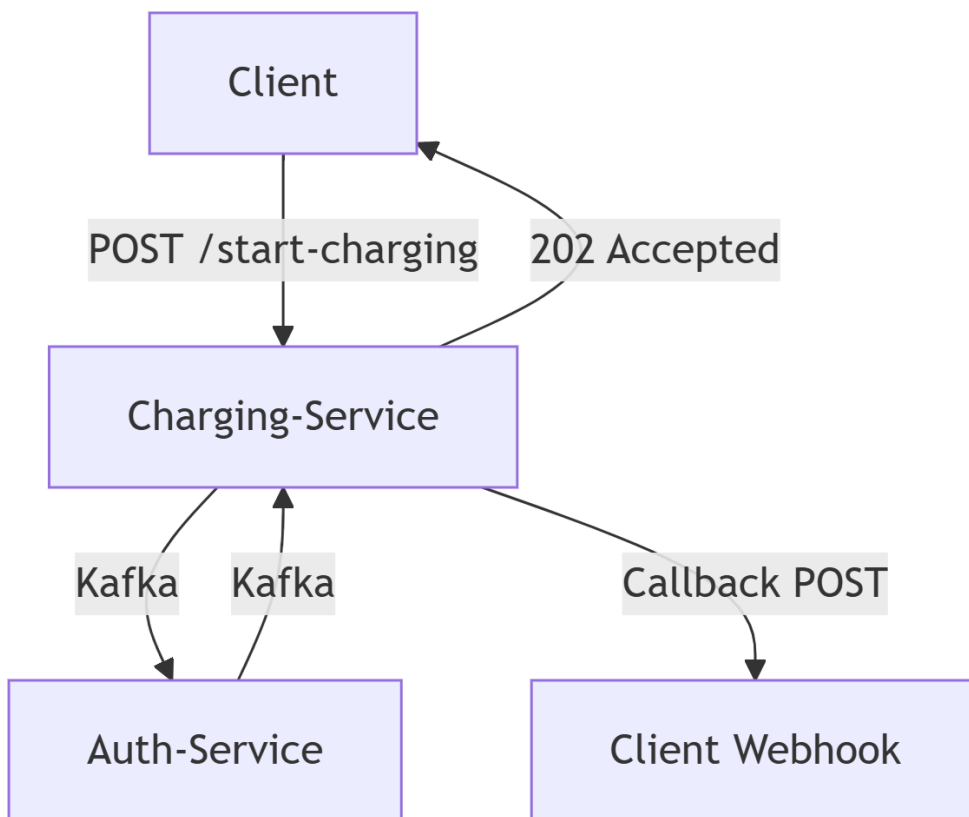


1. Overview

A simple Kotlin, Springboot microservice for managing electric vehicle charging sessions with:

- Asynchronous request processing
- Kafka-based event-driven architecture
- REST API with callback notifications

2. System Architecture



Components

1. **Charge-Service (8082)** - Client request the first request(producer) and consumes the later validated/ authorized response from the auth-service and send the response to callback service.
2. **Auth-Service (8084)** – Consume the request and validate the charge-request (Internal ACL service). Later again send the validated response back to the Auth-Service as a producer.
3. **Kafka (9092)** – Act as a message brocker in between asynchronous commiunication between Charge-Service and Auth-Service
4. **H2 database** – Simple in memory database to persist the result and request.\

3. API Specification

Start Charging Session

Endpoint: POST /api/v1/charging/start-charging

Request: Valid

```
{  
  "stationId": "123e4567-e89b-12d3-a456-426614174000",  
  "driverToken": "validDriverToken-new-123",  
  "callbackUrl": "http://localhost:8082/api/v1/callback/get-callback"  
}
```

Response (202 Accepted):

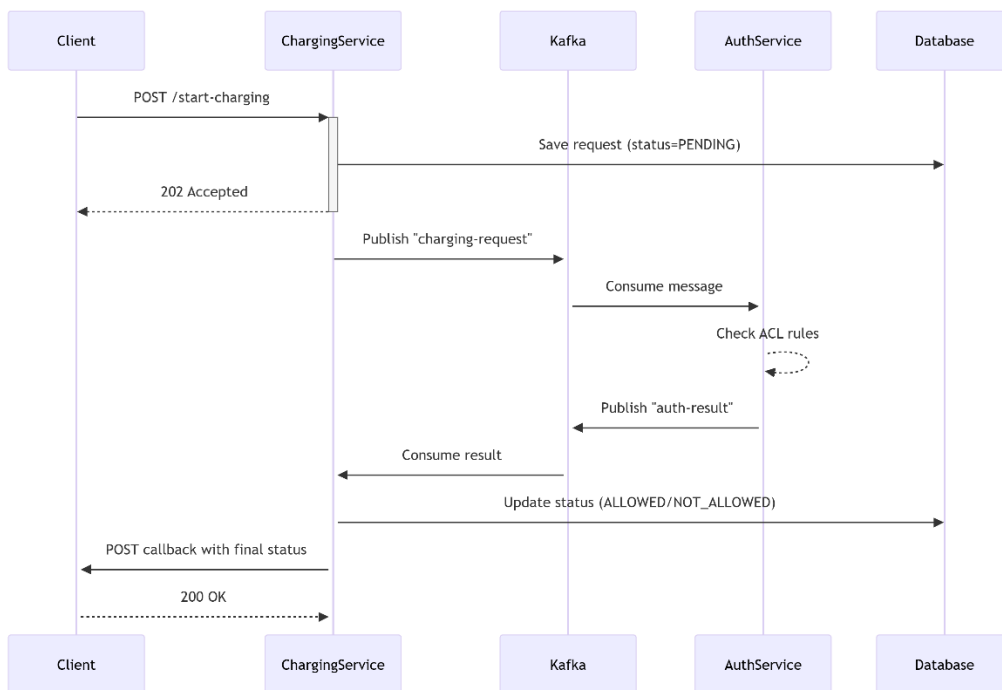
```
{  
  "status": "accepted",  
  "message": "Request is being processed. The result will send to callback Url"  
}
```

You can able to test the api via swagger (/swagger-ui/index.html#/charging-controller/startCharging) or via Postman.

Curl

```
curl -X 'POST' \  
  'http://localhost:8082/api/v1/charging/start-charging' \  
  -H 'accept: */*' \  
  -H 'Content-Type: application/json' \  
  -d '{  
    "stationId": "123e4567-e89b-12d3-a456-426614174000",  
    "driverToken": "validDriverToken-new-1233",  
    "callbackUrl": "http://localhost:8082/api/v1/callback/get-callback"  
  }'
```

4. WorkFlow Diagram



4. Key Features

1. Asynchronous Processing

- Immediate 202 response
- Results delivered via webhook

2. Fault Tolerance

- Default "UNKNOWN" status on timeouts
- Retry mechanism for failed callbacks

3. Validation

- Driver Token: Uppercase letters (A-Z), Lowercase letters (a-z) Digits (0-9), Hyphen (-), period (.), underscore (_), and tilde (~).

@field:NotBlank

@field:Size(min = 20, max = 80)

@field:Pattern(regex = "^[a-zA-Z0-9\\-\\._~]+\$")

val driverToken: String

- StationId : Should be a UUID

@field:NotNull

val stationId: UUID,

- CallbackURL : Should be a valid Http or Https endpoint

@field:NotBlank

@field:Pattern(regex = "https?:/.+")

val callbackUrl: String

5. Solution Design

Problem Solved

- Prevents service overload during peak demand
- Decouples authorization checks from API layer
- Provides real-time notifications without polling

6. Deployment Guide

Prerequisites

- Docker 20+
- JDK 17

You can download the repository from <https://github.com/charithellawala/sample-async-project> and run "docker-compose up --build".

Steps

1. Build and run:

```
docker-compose up --build
```

2. Verify services:

```
docker-compose ps
```

3. Test API:

```
curl -X 'POST' \
  'http://localhost:8082/api/v1/charging/start-charging' \
  -H 'accept: */*' \
  -H 'Content-Type: application/json' \
  -d '{
    "stationId": "123e4567-e89b-12d3-a456-426614174000",
    "driverToken": "validDriverToken-new-1233",
    "callbackUrl": "http://localhost:8082/api/v1/callback/get-callback"
  }'
```

7. Technical Choices

Kafka: Guaranteed message delivery

H2: Lightweight and simple persistence for the sample solution

Spring Boot: Rapid Development with Kotlin

8. Limitations of the Solution.

Kafka Related Limitations :

- No built-in retries for failed callbacks: Lost callbacks if client endpoint is down
- Retry Mechanism still to be designed/implemented for failure scenarios.
- Single Kafka Cluster.

Database Limitations :

- Lack of historical data.
- Bottleneck for high-volume stations.
- Data loss on abrupt shutdown

Architectural Limitations :

- Increased latency during peaks
- Cascading failures if Kafka/AuthService is down
- Hardcoded timeouts

Security Limitations :

- Potential data leakage in Kafka
- Callback URL spooling/ malicious redirects
- No ratelimits or Gateway

Operational Limitations :

- Hard to debug cross-service flows/ No standard tracing.
- Manual scaling

Business Logic Limitations :

- Duplicate charges possible.
- No proper ACL service.
- No pricing model available/ mapped

